

Introduction to Reinforcement Learning II

ML on quantum and classical data

Gorka Muñoz-Gil (**ICFO**)


Master in Photonics (UPC-UAB-UB-ICFO), 2018-2019

Markov decision process

MAB: the action was done over the same *invariant* system. What happens when our actions affect the system?

→ We will need to learn how to choose different actions in different situations! Namely: $q_*(a) \rightarrow q_*(s, a)$

Before that, what is a **MDP**?




Elements of MPD

Formally, an MDP is an Markovian process ('no memory') define by a 5 elements:

- Finite set of states S .
- Finite set of actions A .
- Set of probabilities $P_a(s, s')$
- Set of rewards $R_a(s, s')$
- γ discount factor, which represents the difference in importance between future and past.

Markov decision process

In terms of RL, we treat with systems like



To keep track of learning, we will use the tuple $(S_t, A_t, S_{t+1}, R_{t+1})$.

We will consider *finite* MDP: finite set of states, actions and rewards.

Markov decision process

The MDP is completely characterized by the probability

$$p(s', r|s, a) = \Pr \{S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a\}$$

See that $\sum_{s' \in \mathcal{S}} \sum_{r \in \mathcal{R}} p(s', r|s, a) = 1$. From here one can set interesting probabilistic relations such as the *expected reward for state-action pairs*,

$$r(s, a) = \mathbb{E} [R_t | S_{t-1} = s, A_{t-1} = a] = \sum_{s' \in \mathcal{S}} r \sum_{r \in \mathcal{R}} p(s', r|s, a)$$

In general, MDPs allow for a simple and elegant representation of our real world problem. The only requirement is to draw the correct agent-environment system.

Goals and Rewards

In MDPs, the goal is to maximize the expected value of the cumulative sum of received rewards.

Rewards can be either positive, negative (usually referred as penalization) or zero.

Example

Positive reward: when making a robot learn to walk, reward is proportional to robot's forward motion.

Penalization: when making a robot learn to escape a maze as faster as possible, each step in maze give reward -1.

Returns and Episodes

As in the MAB, we want to maximize the expected return

$$G_t = R_{t+1} + R_{t+2} + \dots + R_T.$$

We consider the agent-environment interact broken into *episodes* of length T . Normally they all start in same initial state.

Additionally, we may want to choose A_t to maximize the expected *discounted return*:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = R_{t+1} + \gamma G_{t+1} = \sum_{k=0}^{T-1} \gamma^k R_{t+k+1}$$

Markov decision process

Policies and value functions

Policy $\pi(a|s)$: mapping from states to probabilities of selecting each possible action. Is the 'brain' of the agent.

Value function $v_\pi(s)$: is the expected return when starting in state s and follow policy π ,

$$v_\pi(s) = \mathbb{E}[G_t | S_t = s] = \mathbb{E}_\pi \left[\sum_{k=0}^{T-1} \gamma^k R_{t+k+1} \middle| S_t = s \right]$$

The value of taking action a in state s , the action-value function, under policy π is defined as

$$q_\pi(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a]$$

Markov decision process

Optimal policies and value functions

Policy π is better than policy π' (also written as $\pi > \pi'$) if and only if $v_\pi(s) > v_{\pi'} \forall s$. We denote the optimal policy as π_* . We then have

$$\boxed{q_* = \max_{\pi} q_{\pi}(s, a)} \quad \boxed{v_* = \max_{\pi} v_{\pi}(s)}$$

Both functions follow the *Bellman optimality equation*:

$$v_*(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]$$
$$q_*(s, a) = \sum_{s', r} p(s', r | s, a) [r + \gamma \max_{a'} q_*(s', a')]$$

Once here, we have now all the tools needed to create our own policy!